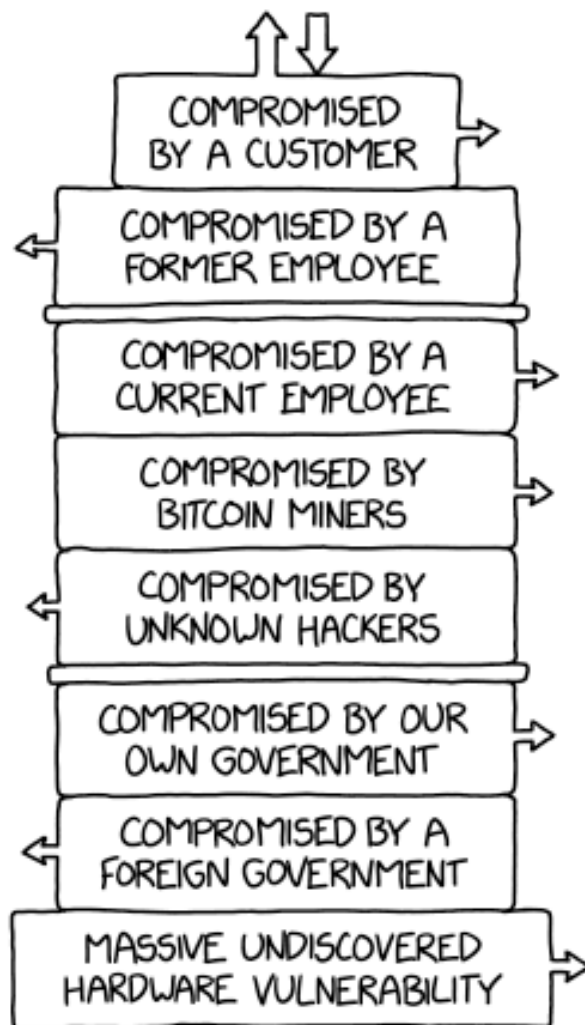


## THE MODERN TECH STACK



# Instruction Set

- Different computers have different instruction sets
  - But with many aspects in common
- You might think that the languages of computers would be as diverse as those of people, but in reality computer languages are quite similar, more like regional dialects than like independent languages.
- The chosen instruction set comes from MIPS Technologies, and is an elegant example of the instruction sets designed since the 1980s.

# Instruction Set

- ARMv7 is similar to MIPS. More than 9 billion chips with ARM processors were manufactured in 2011
- Intel x86, which powers both the PC and the cloud of the PostPC Era.
- ARMv8, which extends the address size of the ARMv7 from 32 bits to 64 bits. Ironically, this 2013 instruction set is closer to MIPS than it is to ARMv7

# Stored-program concept

- Storage of instructions in computer memory to enable it to perform a variety of tasks in sequence or intermittently
- Created by Jon von Neumann in the late 1940s.
- Manchester Mark 1 – proposed and implemented in England
- EDVAC -- proposed and implemented in the US.

# Operations of the Computer Hardware

- The MIPS assembly language notation

`add a, b, c`

instructs a computer to add the two variables `b` and `c` and to put their sum in `a`

- MIPS arithmetic instruction performs only one operation and must always have exactly three variables.

# MIPS Example:

- For example, suppose we want to place the sum of four variables `b`, `c`, `d`, and `e` into variable `a`.

```
add a, b, c # The sum of b and c is placed in a
add a, a, d # The sum of b, c, and d is now in a
add a, a, e # The sum of b, c, d, and e is now in a
```

every newline is a new instruction. No ;

"#" is how to comment

# Example: C $\rightarrow$ MIPS

- This segment of a C program contains the five variables `a`, `b`, `c`, `d`, and `e`

```
a = b + c;
```

```
d = a - e;
```

- A MIPS instruction operates on two source operands and places the result in one destination operand.



# MIPs Instruction Goals:

- Showing how it is represented in hardware
- The relationship between high-level programming languages and this more primitive one
  - Examples in C and JAVA
- See the impact of programming languages and compiler optimization on performance
- writing programs in the language of the computer and running them on the simulator

# Register Operands

- Arithmetic instructions use register operands
- MIPS has a  $32 \times 32$ -bit register file
  - Use for frequently accessed data
  - Numbered 0 to 31
  - 32-bit data called a “word”
- Assembler names
  - \$t0, \$t1, ..., \$t9 for temporary values
  - \$s0, \$s1, ..., \$s7 for saved variables
- Design Principle 2: Smaller is faster

# Example: C $\rightarrow$ MIPS

- A somewhat complex statement contains the five variables `f`, `g`, `h`, `i`, and `j`:

```
f = (g + h) - (i + j);
```

- Need to use temporary variables
- Break it into 3 different MIPS instructions

# Register Operand Example

- C code:

```
f = (g + h) - (i + j);
```

ASSUME: f, g, h, i, and j are assigned to the registers \$s0, \$s1, \$s2, \$s3, and \$s4, respectively. Compiled MIPS code:

```
add $t0, $s1, $s2 # register $t0 contains g + h
add $t1, $s3, $s4 # register $t1 contains i + j
sub $s0, $t0, $t1 # f gets $t0 - $t1, which is (g + h) - (i + j)
```

# Data Types

- Instructions are all 32 bits
- byte(8 bits), halfword (2 bytes), word (4 bytes)
- a character requires 1 byte of storage
- an integer requires 1 word (4 bytes) of storage

# Immediate Operands

- Constant data specified in an instruction

```
addi $s3, $s3, 4 # add immediate $s3 = $s3 + 4
```

- No subtract immediate instruction

- Just use a negative constant

```
addi $s2, $s1, -1
```

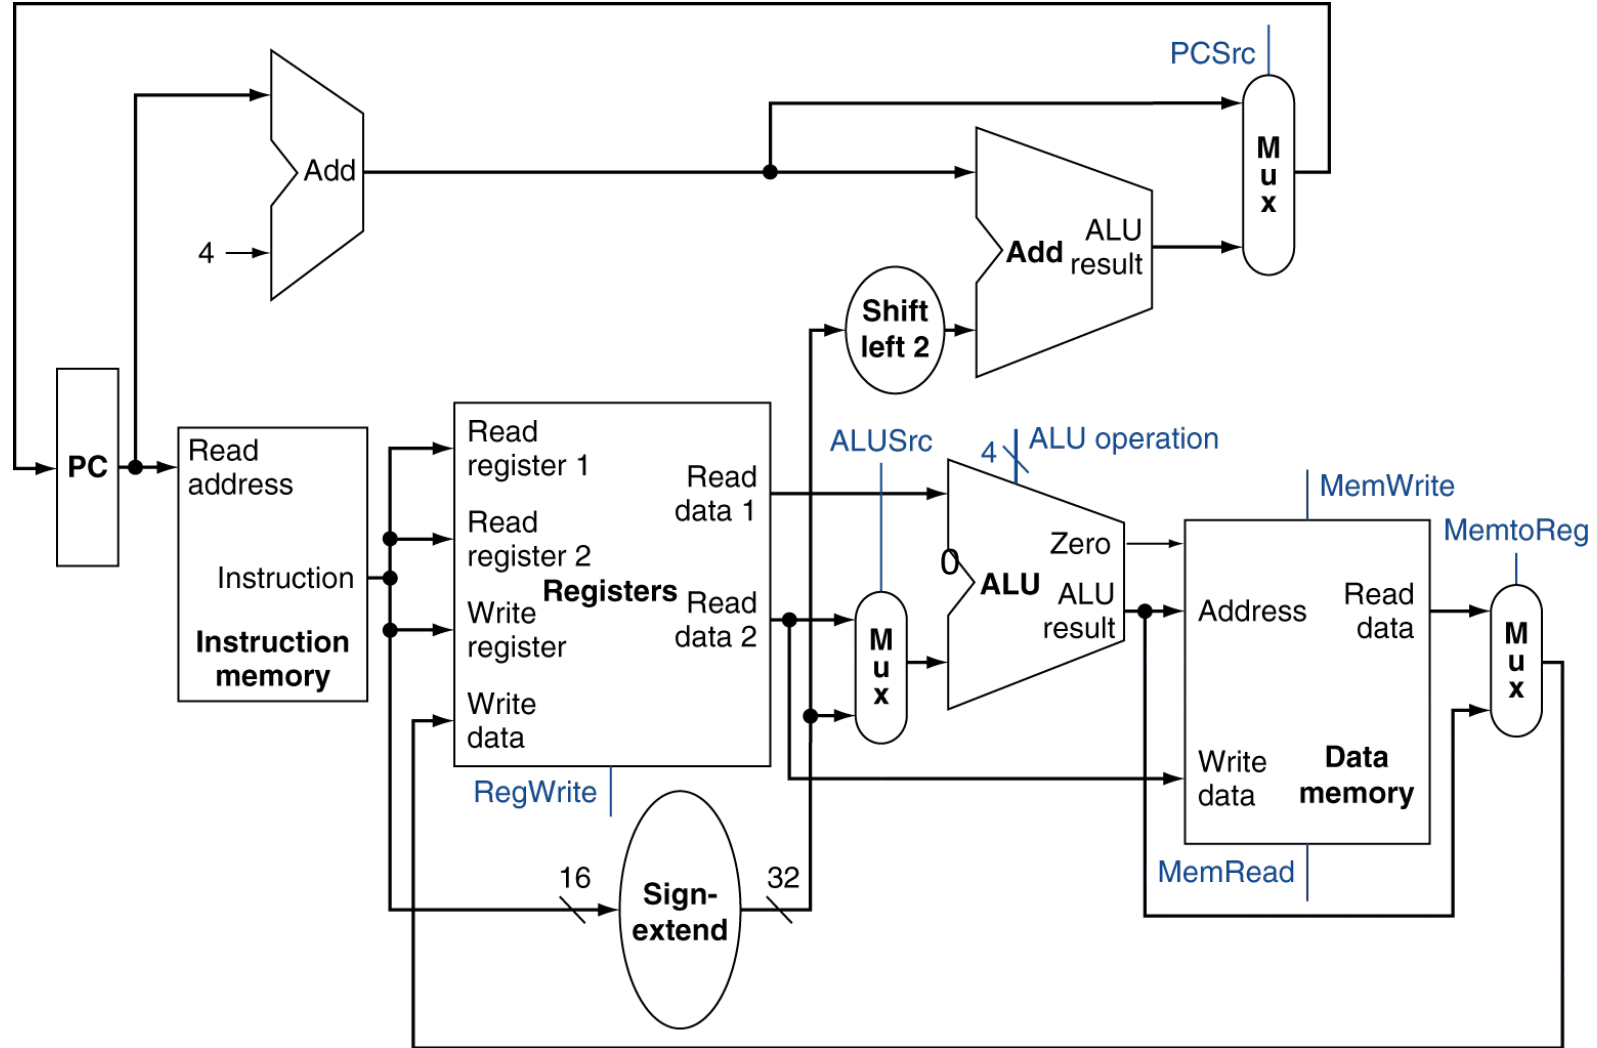
- Design Principle 3: Make the common case fast

- Small constants are common
- Immediate operand avoids a load instruction

# The Constant Zero Register

- MIPS register 0 (\$zero) is the constant 0
  - Cannot be overwritten
- Useful for common operations
  - E.g., move between registers  
add \$t2, \$s1, \$zero

# MIPS Architecture





# Memory Operand

To apply arithmetic operations

- Load values from memory into registers
- Store result from register to memory

# Memory Operand

---

lw    \$t0, 32(\$s3)    # load word  
add \$s1, \$s2, \$t0

offset

base register

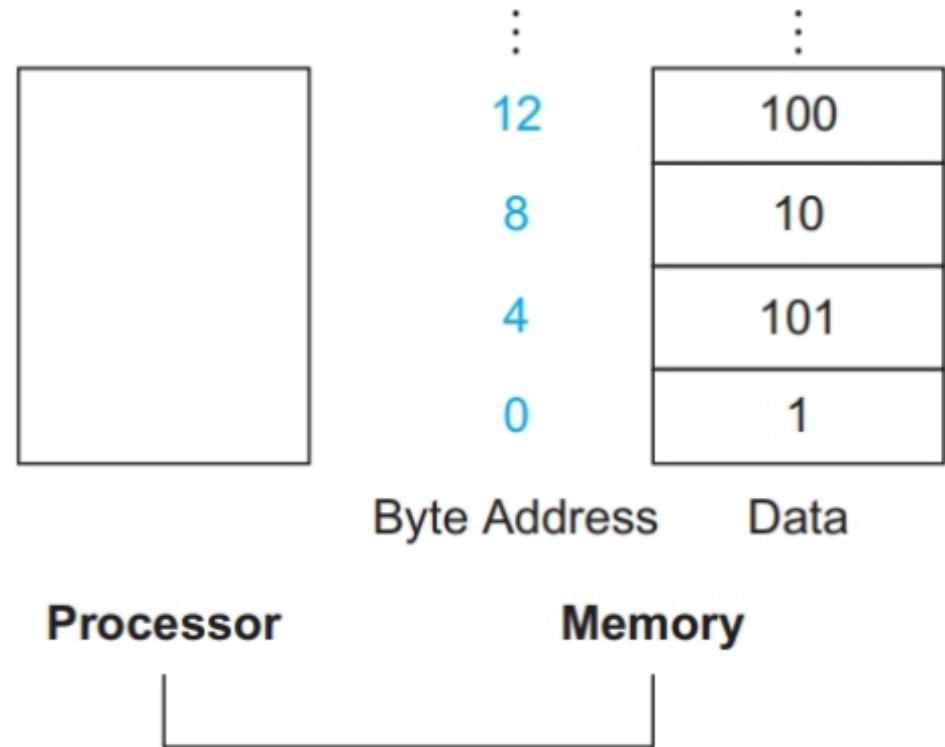
## Memory Operand

`g = h + A[8];`

`g` in `$s1`, `h` in `$s2`, base address of `A` in `$s3`

- Index 8 requires offset of 32

```
lw $t0, 32($s3) # load word
add $s1, $s2, $t0
```



# Memory Operand Example

Assuming  $h \rightarrow s2$  and Base  
address of  $A \rightarrow s3$

  
lw \$t0, 32(\$s3) # Temporary reg \$t0 gets A[8]  
add \$t0, \$s2, \$t0 # Temporary reg \$t0 gets h+A[8]

  
sw \$t0, 48(\$s3) # Stores h+A[8] back into A[12]

$$A[12] = h + A[8]$$

Register Number	Alternative Name	Description
0	zero	the value 0
1	\$at	(assembler temporary) reserved by the assembler
2-3	\$v0 - \$v1	(values) from expression evaluation and function results
4-7	\$a0 - \$a3	(arguments) First four parameters for subroutine. Not preserved across procedure calls
8-15	\$t0 - \$t7	(temporaries) Caller saved if needed. Subroutines can use w/out saving. Not preserved across procedure calls
16-23	\$s0 - \$s7	(saved values) - Callee saved. A subroutine using one of these must save original and restore it before exiting. Preserved across procedure calls
24-25	\$t8 - \$t9	(temporaries) Caller saved if needed. Subroutines can use w/out saving. These are in addition to \$t0 - \$t7 above. Not preserved across procedure calls.
26-27	\$k0 - \$k1	reserved for use by the interrupt/trap handler
28	\$gp	global pointer. Points to the middle of the 64K block of memory in the static data segment.
29	\$sp	stack pointer Points to last location on the stack.
30	\$s8/\$fp	saved value / frame pointer Preserved across procedure calls
31	\$ra	return address

# MIPS Instruction Computer Representation

- Instructions are kept in the computer as a series of high and low electronic signals
- May be represented as numbers
- Each piece of an instruction can be considered as an individual number
- Placing these numbers side by side forms the instruction.

# MIPS Instruction Computer Representation

add \$t0, \$s1, \$s2

Decimal:

0	17	18	8	0	32
---	----	----	---	---	----

  
Field

Binary:

000000	10001	10010	01000	00000	100000
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

Instruction Format – Machine Language



**op:** Basic operation of the instruction, traditionally called the opcode.

**rs:** The first register source operand

**rt:** The second register source operand.

MIPS fields Defined





**rd:** The register destination operand. It gets the result of the operation

**shamt:** Shift amount

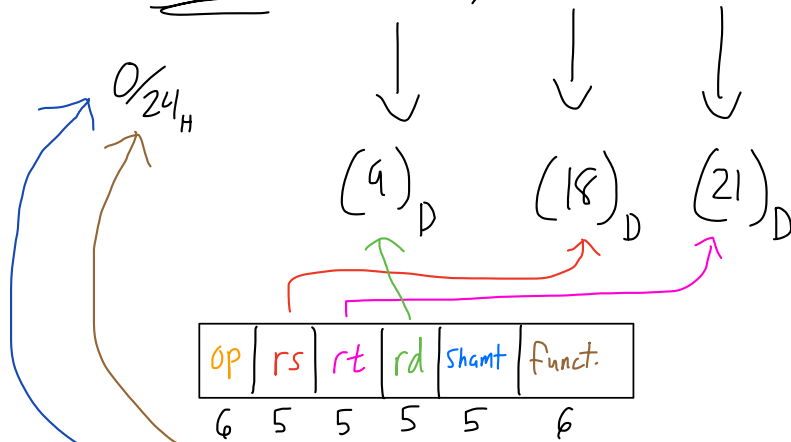
**funct:** Function. This field, often called the function code, selects the specific variant of the operation in the op field

MIPS fields Defined

# MIPS Instructions

- keep all instructions the same length
- Require different kinds of instruction formats for different kinds of instructions.
- R-format: r for registers (R-type)
- I-format: i for immediate (I-type)
- The formats are distinguished by the values in the first field each format is assigned a distinct set of values in the first field (op)
- the hardware knows whether to treat the last half of the instruction as three fields (R-type) or as a single field (I-type)

Ex. AND \$t1, \$s2, \$s5



NAME	NUMBER
\$zero	0
\$at	1
\$v0-\$v1	2-3
\$a0-\$a3	4-7
\$t0-\$t7	8-15
\$s0-\$s7	16-23
\$t8-\$t9	24-25
\$k0-\$k1	26-27
\$gp	28
\$sp	29
\$fp	30
\$ra	31

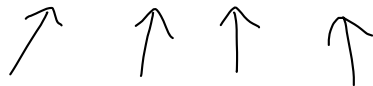
0010 = 2

0100 = 4

But we only have 6 Bits  
so we chop the leading zeros on 2

2<sub>H</sub> = 100100

000000 10010 10101 01001 00000 100100



make sure to pad the left to make sure it's 5 bits

BTW: The largest Num you can  
have in the function section  
is 3 F because it's 6 bits

11 111

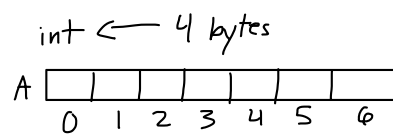
Slides 17-20

Ex.

lw \$t0, 32 (\$s3)

The parentheses tell me that I want the  
value at that address and not the address

Base memory address



A[7]

I want to  
access A[0]

# assume base address A is in \$s2

lw \$t1, 0 (\$s2)

Base address

This loads the value of s2 into t1 and not the address

34 mins

Ex.

Take that yellow line of code and get the mips instructions for it

$A[300] = h + A[300];$

$\$t1 \rightarrow$  Base of A     $\$s2 \rightarrow h$   
          

$lw \$t0, 1200(\$t1) \quad \# \$t0 = A[300]$   
 $add \$t0, \$s2, \$t0 \quad \# h + A[300]$   
 $sw \$t0, 1200(\$t1) \quad \# A[300] = \$t0$

now I'm writing out the Binary sequence for this

This is a new question

op = 23<sub>hex</sub>

rs = 9

rt = 8

Const = 1200

| Load Word    lw    I     $R[rt] = M[R[rs] + \text{SignExtImm}]$     (2) 23<sub>hex</sub>

### I-format Instructions

op	rs	rt	constant or address
6 bits	5 bits	5 bits	16 bits

010011    01001    01000    0000010010110000  
op            rs            rt            const

# R-format Example

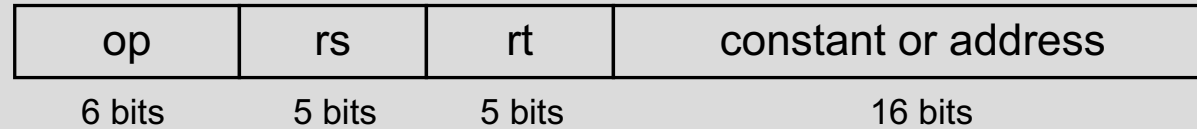
op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

add \$t0, \$s1, \$s2

special	\$s1	\$s2	\$t0	0	add
0	17	18	8	0	32
000000	10001	10010	01000	00000	100000

$00000010001100100100000000100000_2 = 02324020_{16}$

# I-format Instructions



- Immediate arithmetic and load/store instructions
  - rt: destination or source register number
  - Constant:  $-2^{15}$  to  $+2^{15} - 1$
  - Address: offset added to base address in rs
- *Design Principle 4: Good design demands good compromises*
  - Different formats complicate decoding, but allow 32-bit instructions uniformly
  - Keep formats as similar as possible

Instruction	Format	op	rs	rt	rd	shamt	funct	address
add	R	0	reg	reg	reg	0	32 <sub>ten</sub>	n.a.
sub (subtract)	R	0	reg	reg	reg	0	34 <sub>ten</sub>	n.a.
add immediate	I	8 <sub>ten</sub>	reg	reg	n.a.	n.a.	n.a.	constant
lw (load word)	I	35 <sub>ten</sub>	reg	reg	n.a.	n.a.	n.a.	address
sw (store word)	I	43 <sub>ten</sub>	reg	reg	n.a.	n.a.	n.a.	address

# Example:

We can now take an example all the way from what the programmer writes to what the computer executes. If `$t1` has the base of the array `A` and `$s2` corresponds to `h`, the assignment statement

```
A[300] = h + A[300];
```

is compiled into

```
lw    $t0,1200($t1) # Temporary reg $t0 gets A[300]
add   $t0,$s2,$t0   # Temporary reg $t0 gets h + A[300]
sw    $t0,1200($t1) # Stores h + A[300] back into A[300]
```



# Example: Machine Language Decimal Solution

Op	rs	rt	rd	address/ shamt	funct
35	9	8	1200		
0	18	8	8	0	32
43	9	8	1200		

# Example: Machine Language Binary Solution

Op	rs	rt	rd	address/ shamt	funct
100011	01001	01000	0000 0100 1011 0000		
000000	10010	01000	01000	00000	100000
101011	01001	01000	0000 0100 1011 0000		

# Sign Extension

- Representing a number using more bits
  - Preserve the numeric value
- When loading a 32-bit word into a 32-bit register, signed and unsigned loads are identical.
- MIPS does offer two flavors of byte loads:
  - ***load byte (lb)***: treats the byte as a signed number and thus sign-extends to fill the 24 left-most bits of the register,
  - ***load byte unsigned (lbu)***: works with unsigned integers.

Examples: 8-bit to 16-bit

+2: 0000 0010 => 0000 0000 0000 0010

-2: 1111 1110 => 1111 1111 1111 1110

# Logical Operators

Logical operations	C operators	Java operators	MIPS instructions
Shift left	<<	<<	sll
Shift right	>>	>>>	srl
Bit-by-bit AND	&	&	and, andi
Bit-by-bit OR			or, ori
Bit-by-bit NOT	~	~	nor

# Shift Operations

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

- shamt: how many positions to shift
- Shift left logical
  - Shift left and fill with 0 bits
  - $sll$  by  $i$  bits multiplies by  $2^i$
- Shift right logical
  - Shift right and fill with 0 bits
  - $srl$  by  $i$  bits divides by  $2^i$  (unsigned only)

# Shift

---

```
sll $t2,$s0,4  # reg $t2 = reg $s0 << 4 bits
```

op	rs	rt	rd	shamt	funct
0	0	16	10	4	0

0000 0000 0000 0000 0000 0000 0000 1001<sub>two</sub> = 9<sub>ten</sub>

0000 0000 0000 0000 0000 0000 1001 0000<sub>two</sub> = 144<sub>ten</sub>

# NOT Operations

- Useful to invert bits in a word
  - Change 0 to 1, and 1 to 0
- MIPS has NOR 3-operand instruction
  - $a \text{ NOR } b == \text{NOT} (a \text{ OR } b)$

`nor $t0, $t1, $zero` ←

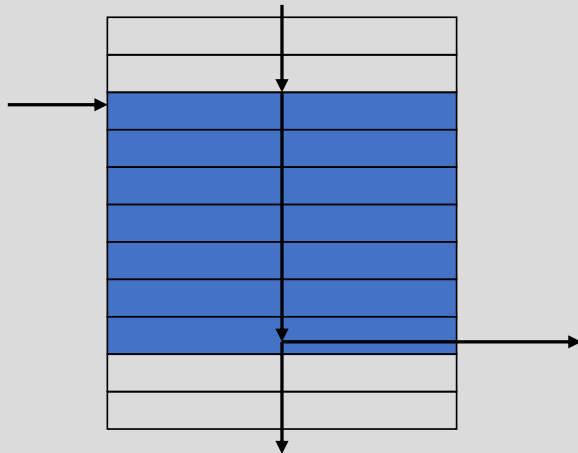
Register 0: always  
read as zero

\$t1 0000 0000 0000 0000 0011 1100 0000 0000

\$t0 1111 1111 1111 1111 1100 0011 1111 1111

# Basic Blocks

- A basic block is a sequence of instructions with
  - No embedded branches (except at end)
  - No branch targets (except at beginning)



- A compiler identifies basic blocks for optimization
- An advanced processor can accelerate execution of basic blocks



# Labels

- In MIPS assembly, a label is simply a string used to name a location in memory.
- A label may refer to the location of a data value (variable) or of an instruction.
- Think of a label as representing an address
- Labels are terminated by a colon character

# Complete MIPS Example

```
.data
N:      .word 10

.text

main:
    lw    $t0, N        # $t0 <-- Mem[N] (10)
    la    $t1, N        # $t1 <-- N      (address)
    . . .
exit:   li    $v0, 10
        syscall
```

# Unconditional Branch Instructions

- `j Label           # PC = Label`
- `b Label           # PC = Label`
- `jr $ra           # PC = $ra`
- These help with constructing loops

# Making Decisions

- MIPS assembly language includes two decision-making instructions, similar to an if statement with a go to.

```
beq register1, register2, L1
```

- branch if equal – go to L1 if register1 is equal to register2

```
bne register1, register2, L1
```

- branch if not equal

# Example: Decisions

---

In the following code segment, *f*, *g*, *h*, *i*, and *j* are variables. If the five variables *f* through *j* correspond to the five registers \$s0 through \$s4, what is the compiled MIPS code for this C *if* statement?

```
if (i == j) f = g + h; else f = g - h;
```

# Solution

```
bne $s3,$s4,Else      # go to Else if i ≠ j
add $s0,$s1,$s2        # f = g + h (skipped if i ≠ j)

j Exit      # go to Exit

Else: sub $s0,$s1,$s2 # f = g - h (skipped if i = j)
Exit:
```

# Example: Loops

---

Here is a traditional loop in C:

```
while (save[i] == k)
    i += 1;
```

Assume that `i` and `k` correspond to registers `$s3` and `$s5` and the base of the array `save` is in `$s6`. What is the MIPS assembly code corresponding to this C segment?

# Loop

1. load save[i] into a temporary register.
  - Base Address of save
  - Multiply i to get the byte address of the index (by 4)
2. Loop test
  - Choose to use the bne or beq
3. Make branch labels for each portion of the loop test



## Solution: loops

```
Loop: sll $t1,$s3,2      # Temp reg $t1 = i * 4
      add $t1,$t1,$s6    # $t1 = address of save[i]
      lw $t0,0($t1)      # Temp reg $t0 = save[i]
      bne $t0,$s5, Exit  # go to Exit if save[i] ≠ k
      addi $s3,$s3,1     # i = i + 1
      j Loop             # go to Loop
```

Exit:

# set on less than

```
slt    $t0, $s3, $s4    # $t0 = 1 if $s3 < $s4
```

```
slti   $t0, $s2, 10     # $t0 = 1 if $s2 < 10
```

Note: useful for for loops

**NOTE:** MIPS compilers use the `slt`, `slti`, `beq`, `bne`, and the fixed value of 0 (always available by reading register `$zero`) to create all relative conditions: equal, not equal, less than, less than or equal, greater than, greater than or eq

# Example Switch Statements

- The simplest way to implement switch is via a sequence of conditional tests, turning the switch statement into a chain of if-then-else statements.

What is the MIPS assembly code assuming f-k correspond to registers \$s0-\$s5 and \$t2 contains 4 and \$t4 contains base address of JumpTable?

```
switch (k) {  
    case 0: f=i+j;break;  
    case 1: f=g+h;break;  
    case 2: f=g-h;break;  
    case 3: f=i-j;break;  
}
```

```

slt  $t3,$s5,$zero    # test if k=4
add  $t1,$s5,$s5       #$t1 =2*k
add  $t1,$t1,$t1       #$t1 =4*k
add  $t1,$t1,$t4       #$t1=address of JumpTable[k]
lw   $t0,0($t1)        #$t0=JumpTable[k]
jr   $t0               #jump based on register $t0
L0:   add $s0,$s3,$s4
      j Exit
L1:   add $s0,$s1,$s2
      j Exit
L2:   sub $s0,$s1,$s2
      j Exit
L3:   sub $s0,$s3,$s4

Exit:

```

# Procedures in MIPS

function is one tool programmers use to structure programs, both to make them easier to understand and to allow code to be reused.

1. Put parameters in a place where the procedure can access them.
2. Transfer control to the procedure.
3. Acquire the storage resources needed for the procedure.
4. Perform the desired task.
5. Put the result value in a place where the calling program can access it.
6. Return control to the point of origin, since a procedure can be called from several points in a program.

# Procedures in MIPS

- `$a0–$a3`: four argument registers in which to pass parameters
- `$v0–$v1`: two value registers in which to return values
- `$ra`: one return address register to return to the point of origin
- **jump-and-link instruction (jal)**: it jumps to an address and simultaneously saves the address of the following instruction in register `$ra` (register 31)

# Stack Pointer

- Suppose a compiler needs more registers for a procedure than the four argument and two return value registers.
- Any registers needed by the caller must be restored to the values that they contained before the procedure was invoked.
- The ideal data structure for spilling registers is a *stack*—a last-in-first-out.
- Stack pointer is adjusted by one word for each register that is saved or restored. MIPS software reserves register 29 for the stack pointer, giving it the obvious name \$sp.

## Procedure Example

The parameter variables g, h, i, and j correspond to the argument registers \$a0, \$a1, \$a2, and \$a3, and f corresponds to \$s0.

```
int leaf_example (int g, int h, int i, int j)
{
    int f;

    f = (g + h) - (i + j);
    return f;
}
```



# Leaf Procedure Example

- MIPS code:

leaf_example:		
addi	\$sp, \$sp, -4	Save \$s0 on stack
sw	\$s0, 0(\$sp)	
add	\$t0, \$a0, \$a1	Procedure body
add	\$t1, \$a2, \$a3	
sub	\$s0, \$t0, \$t1	
add	\$v0, \$s0, \$zero	Result
lw	\$s0, 0(\$sp)	Restore \$s0
addi	\$sp, \$sp, 4	
jr	\$ra	Return

High address

\$sp →

Low address

\$sp →

Contents of register \$t1

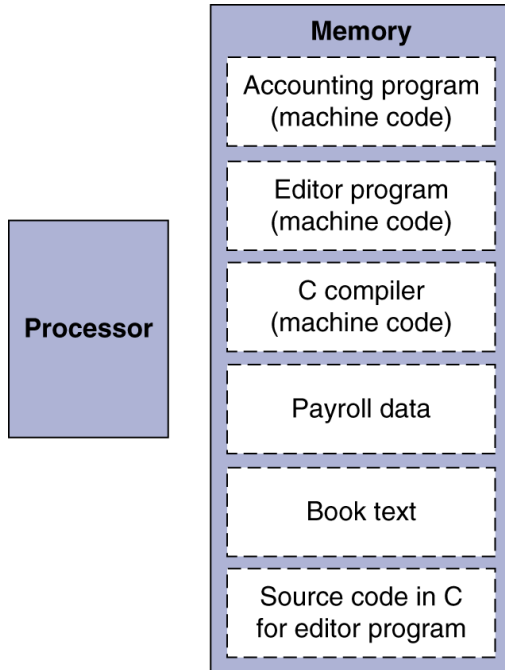
Contents of register \$t0

Contents of register \$s0

\$sp →

# Stored Program Computers

## The BIG Picture



- Instructions represented in binary, just like data
- Instructions and data stored in memory
- Programs can operate on programs
  - e.g., compilers, linkers, ...
- Binary compatibility allows compiled programs to work on different computers
  - Standardized ISAs

